

Task 2: Values Audit

Item	Selection
Regulated entity	Citibank, N.A. / Citi Singapore
Domain	AML/KYC, suspicious transaction monitoring, and sanctions-related alert calibration
Jurisdictions	United States and Singapore
Hypothetical RegTech company	AZ RegTech

1. Company mission statement

AZ RegTech is a hypothetical early-stage compliance technology company that helps global banks calibrate AML alerts across jurisdictions. Its mission is to help financial institutions detect suspicious activity consistently, explainably, and proportionately, while avoiding a compliance process that is reduced to box-ticking or indiscriminate alert generation.

The company is built around three core values: substance over paperwork, accountable automation, and jurisdictional humility. Substance over paperwork means that the tool should improve the quality of risk detection rather than merely produce a larger audit trail. Accountable automation means that software can prioritise, score, and explain alerts, but it should not replace the professional judgement required for suspicious transaction reporting. Jurisdictional humility means that the tool should not pretend that US and Singapore AML requirements are identical just because both jurisdictions use risk-based language.

AZ is assumed to be a Series A-stage company with around 35 employees, including engineers, data analysts, AML subject matter experts, and implementation specialists. Its initial customers are large banks with cross-border operations. Its revenue model is annual enterprise subscription pricing, supported by implementation fees for rule configuration, data integration, and model validation support. The company is not trying to serve every country immediately. Its first commercial aspiration is to become credible in a small number of high-value financial centres, beginning with the United States, Singapore, Hong Kong, and the United Kingdom.

The company is good at translating regulatory expectations into configurable monitoring logic. It is not a law firm and does not provide legal advice. Its advantage is operational: it helps compliance teams understand how changes in thresholds, customer risk factors, sanctions scores, and transaction typologies affect alert volume, false-positive workload, and escalation quality.

2. Whose perspective does the tool serve?

The primary user served by the tool is the paying client: Citi's Chief Compliance Officer, AML compliance leadership, and financial crime operations teams. These users need a practical way to calibrate alerts across US and Singapore settings without overwhelming investigators or ignoring jurisdiction-specific risk factors. The tool is designed to help them answer questions such as: which transactions should be prioritised, why was an alert generated, which jurisdictional configuration was active, and how much manual review workload will result from a threshold change?

However, the tool cannot be evaluated only from the bank's perspective. It also affects regulators, customers, and the public. Regulators want suspicious activity to be detected, reviewed, and reported in a meaningful way. Customers want their legitimate transactions to proceed without unnecessary delay, account restriction, or reputational harm. The public has an interest in banks not being used to launder criminal proceeds, evade sanctions, or finance terrorism. These interests overlap, but they are not identical.

Stakeholder	Interest	Potential tension
Citi compliance team	Detect reportable activity and manage investigation workload.	May prefer fewer false positives, but excessive filtering can create false negatives.
Regulators	Receive meaningful suspicious transaction reporting and see evidence of effective controls.	May distrust alert reductions if they look like cost-cutting rather than better calibration.
Customers	Avoid unnecessary delays, account reviews, or exits from banking relationships.	Customer convenience may conflict with enhanced due diligence and sanctions review.
Public interest	Prevent money laundering, sanctions evasion, and terrorist financing.	Public safety may justify intrusive review, but overbroad monitoring can unfairly burden legitimate customers.

My design choice is therefore to serve the bank's compliance team directly, while making the tool's outputs explainable enough to be defensible to regulators and constrained enough to reduce avoidable customer harm. I would not market the product as a tool that maximises SAR or STR volume. High volume is not the same as high-quality risk detection.

3. Genuine risk detection or documentation of compliance?

The tool is intended to measure genuine financial crime risk, not merely document compliance. The most important design choice reflecting this value is the inclusion of alert-quality metrics alongside alert-count metrics. A tool that only reports the number of alerts generated can encourage defensive over-reporting. It may look active, but it can

bury investigators under low-value alerts and make it harder to identify truly suspicious behaviour.

For Task 3, the tool will therefore monitor three quantitative outputs: alert rate, false-positive proxy rate, and manual review workload. The false-positive proxy will be estimated using synthetic labels for transactions that trigger rule-based alerts but lack stronger suspicious-pattern indicators. This is not a perfect measure of true false positives, but it forces the design to ask whether extra alerts are likely to add risk insight or merely create operational noise.

A second design choice is that the tool will not automatically file a US SAR or Singapore STR. It will produce escalation recommendations and jurisdiction-specific reasons, but final reporting judgement remains with trained compliance staff. This boundary is deliberate. Suspicious transaction reporting requires contextual judgement, legal interpretation, and accountability. Automating the final filing decision would make the product appear more powerful, but it would outsource judgement in a way that is inappropriate for this domain.

The tool is also designed to expose jurisdictional assumptions. In US mode, it will put more weight on sanctions scores, structuring-like patterns, and US nexus indicators. In Singapore mode, it will put more weight on beneficial ownership completeness, cross-border wire information, customer due diligence quality, and STRO review indicators. The point is not to claim that one jurisdiction is stricter than the other. The point is to show that the same transaction can be operationally risky for different reasons under different regulatory configurations.

4. Who bears the cost if the tool gets it wrong?

The cost of error is unevenly distributed. This is why the tool must be designed with explicit failure modes rather than sold as a neutral scoring engine. In AML monitoring, both false positives and false negatives are serious, but they harm different stakeholders in different ways.

Failure mode	Specific harm	Stakeholder most affected
False negative	A genuinely suspicious transaction is not escalated, allowing laundering, sanctions evasion, or terrorist financing risk to pass through the bank.	Public interest, regulators, and Citi.
False positive	A legitimate customer is delayed, investigated, or potentially exited from a banking relationship without sufficient cause.	Customer and Citi's compliance operations team.
Performance drift	Changing transaction patterns make the calibration less reliable, causing alert quality to	Citi compliance team and regulators.

	deteriorate over time.	
Jurisdictional misconfiguration	A Singapore transaction is reviewed with US-weighted logic, or a US transaction is reviewed with Singapore-weighted logic, leading to missed local concerns or unnecessary escalation.	Citi, regulators, and affected customers.
Over-documentation	The tool creates a persuasive audit trail without improving the underlying judgement about suspicious activity.	Regulators, public interest, and ultimately Citi.

The most serious social harm is the false negative, since it can allow criminal or sanctioned activity to continue through the financial system. The most visible individual harm is the false positive, because a legitimate customer may face transaction delays, repeated questioning, or loss of access to banking services. The most dangerous institutional harm is jurisdictional misconfiguration, as the bank may believe it is compliant while applying the wrong logic to the wrong regulatory environment.

My preferred trade-off is not to minimise one error type at all costs. Instead, the tool should make the trade-off visible. If lowering a sanctions threshold increases alert volume by 40 percent but only marginally improves detection of high-risk typologies, senior compliance staff should see that consequence before adopting the threshold. If Singapore-mode settings produce fewer sanctions alerts but more beneficial ownership escalations, the tool should explain that this reflects a different regulatory emphasis rather than a weaker control environment.